

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df = pd.read_csv("cars.csv")
df.head()
```

	make	model	priceUSD	year	condition	mileage(kilometers)
fuel_type \						
0	mazda	2	5500	2008	with mileage	162000.0
petrol						
1	mazda	2	5350	2009	with mileage	120000.0
petrol						
2	mazda	2	7000	2009	with mileage	61000.0
petrol						
3	mazda	2	3300	2003	with mileage	265000.0
diesel						
4	mazda	2	5200	2008	with mileage	97183.0
diesel						

	volume(cm3)	color	transmission	drive_unit	segment
0	1500.0	burgundy	mechanics	front-wheel drive	B
1	1300.0	black	mechanics	front-wheel drive	B
2	1500.0	silver	auto	front-wheel drive	B
3	1400.0	white	mechanics	front-wheel drive	B
4	1400.0	gray	mechanics	front-wheel drive	B

*# to check shape of the dataset*

```
df.shape
```

```
(56244, 12)
```

```
df.dtypes
```

```
make          object
model         object
priceUSD      int64
year          int64
condition     object
mileage(kilometers) float64
fuel_type     object
volume(cm3)   float64
color         object
transmission  object
drive_unit    object
segment       object
dtype: object
```

```

# dropping unnecessary tables from the dataset for the analysis
df.drop(columns = ['model', 'segment'], inplace=True)

# to find the unique values in the columns
df.nunique()

make          96
priceUSD      2970
year          78
condition      3
mileage(kilometers)  8400
fuel_type      3
volume(cm3)    458
color         13
transmission   2
drive_unit     4
dtype: int64

# gives me all kinds of makes/types from the dataset
# it can be done for all the columns
print(df['make'].unique())
print("colors")
print(df['color'].unique())

['mazda' 'mg' 'renault' 'gaz' 'aro' 'rover' 'uaz' 'alfa-romeo' 'audi'
 'oldsmobile' 'saab' 'peugeot' 'chrysler' 'wartburg' 'moskvich'
 'volvo'
 'fiat' 'roewe' 'porsche' 'zaz' 'luaz' 'dacia' 'lada-vaz' 'izh' 'raf'
 'bogdan' 'bmw' 'nissan' 'mercedes-benz' 'mitsubishi' 'toyota' 'chery'
 'gmc' 'hyundai' 'honda' 'ssangyong' 'suzuki' 'opel' 'seat'
 'volkswagen'
 'daihatsu' 'chevrolet' 'geely' 'saturn' 'kia' 'lincoln' 'eksklyuziv'
 'citroen' 'dong-feng' 'pontiac' 'ford' 'subaru' 'bentley' 'faw'
 'cadillac' 'lifan' 'plymouth' 'hafei' 'shanghai-maple' 'mini' 'jeep'
 'skoda' 'mercury' 'changan' 'lexus' 'isuzu' 'aston-martin' 'lancia'
 'great-wall' 'land-rover' 'jaguar' 'buick' 'daewoo' 'vortex'
 'infiniti'
 'byd' 'smart' 'maserati' 'haval' 'acura' 'scion' 'tata' 'datsun'
 'tesla'
 'mclaren' 'ravon' 'trabant' 'proton' 'fso' 'jac' 'asia' 'iran-khodro'
 'zotye' 'tagaz' 'saipa' 'brilliance']
colors
['burgundy' 'black' 'silver' 'white' 'gray' 'blue' 'other' 'purple'
 'red'
 'green' 'brown' 'yellow' 'orange']

```

Since there are many cars, it is difficult to analyze them so i will group them into categories: Luxury European, Mainstream European, Russian/ Eastern European, Asian, American, Speciality, and Other.

```

#categorizing the car make according to the categories
def car_make(make):
    if make in ['mazda', 'mg', 'rover', 'alfa-romeo', 'audi', 'peugeot',
'chrysler', 'bmw', 'aston-martin', 'jaguar', 'land-rover']:
        return 'Luxury European'
    elif make in ['renault', 'dacia', 'citroen', 'volvo', 'fiat',
'opel', 'seat', 'volkswagen', 'citroen', 'skoda', 'mini', 'smart' ]:
        return 'Mainstream European'
    elif make in ['gaz', 'aro', 'lada-vaz', 'izh', 'raf', 'bogdan',
'moskvich', 'uaz', 'luaz', 'wartburg', 'trabant', 'proton', 'fso',
'jac', 'iran-khodro', 'zotye', 'tagaz', 'saipa', 'brilliance']:
        return 'Russian/Eastern European'
    elif make in ['toyota', 'nissan', 'asia', 'mitsubishi', 'chery',
'hyundai', 'honda', 'ssangyong', 'suzuki', 'daihatsu', 'kia',
'changan', 'lexus', 'isuzu', 'great-wall', 'daewoo', 'vortex',
'infiniti', 'byd', 'geely', 'haval', 'acura', 'scion', 'tata',
'datsun', 'ravon', 'proton', 'jac']:
        return 'Asian'
    elif make in ['oldsmobile', 'gmc', 'chrysler', 'plymouth', 'ford',
'cadillac', 'jeep', 'mercury', 'lincoln', 'buick', 'saturn',
'pontiac', 'chevrolet']:
        return 'American'
    elif make in ['porsche', 'bentley', 'maserati', 'tesla',
'mclaren']:
        return 'Specialty'
    else:
        return 'Other'
    # this below functions is used for creating a column (car_segment)
    to display the categories in the datasets
df['make_segment'] = df['make'].apply(car_make)

df.describe()

```

	priceUSD	year	mileage(kilometers)	volume(cm3)
count	56244.000000	56244.000000	5.624400e+04	56197.000000
mean	7415.456440	2003.454840	2.443956e+05	2104.860615
std	8316.959261	8.144247	3.210307e+05	959.201633
min	48.000000	1910.000000	0.000000e+00	500.000000
25%	2350.000000	1998.000000	1.370000e+05	1600.000000
50%	5350.000000	2004.000000	2.285000e+05	1996.000000
75%	9807.500000	2010.000000	3.100000e+05	2300.000000
max	235235.000000	2019.000000	9.999999e+06	20000.000000

```
df.head()
```

	make	priceUSD	year	condition	mileage(kilometers)	fuel_type
0	mazda	5500	2008	with mileage	162000.0	petrol
1	mazda	5350	2009	with mileage	120000.0	petrol


2	mazda	7000	2009	with mileage	61000.0	petrol
3	mazda	3300	2003	with mileage	265000.0	diesel
4	mazda	5200	2008	with mileage	97183.0	diesel

	volume(cm3)	color	transmission	drive	unit	
make_segment						
0	1500.0	burgundy	mechanics	front-wheel	drive	Luxury European
1	1300.0	black	mechanics	front-wheel	drive	Luxury European
2	1500.0	silver	auto	front-wheel	drive	Luxury European
3	1400.0	white	mechanics	front-wheel	drive	Luxury European
4	1400.0	gray	mechanics	front-wheel	drive	Luxury European

```
import matplotlib.pyplot as plt
import seaborn as sns





# Assuming you have a DataFrame 'df' with a column 'make_segment'
plt.figure(figsize=(10, 6))
sns.barplot(x='make_segment', y='make_segment', data=df,
            estimator=len)

plt.title('Distribution of Car Make Segments')
plt.xlabel('Make Segment')
plt.ylabel('Count')
plt.xticks(rotation=90)
plt.show()
```









Categorical variable Distribution

```
# creates empty the 2x2 matrix graphs with empty values
fig, ax = plt.subplots(2,2,figsize=(20,10))
```



```
fig, ax = plt.subplots(2,3,figsize=(30,30))
# plots the graph acc to axis the data from df
sns.countplot(x='condition', data=df, ax=ax[0,0])
sns.countplot(x='fuel_type', data=df, ax=ax[0,1])
sns.countplot(x='transmission', data=df, ax=ax[0,2])
sns.countplot(x='color', data=df, ax=ax[1,0])
sns.countplot(x='drive_unit', data=df, ax=ax[1,1])
sns.countplot(x='make_segment', data=df, ax=ax[1,2])

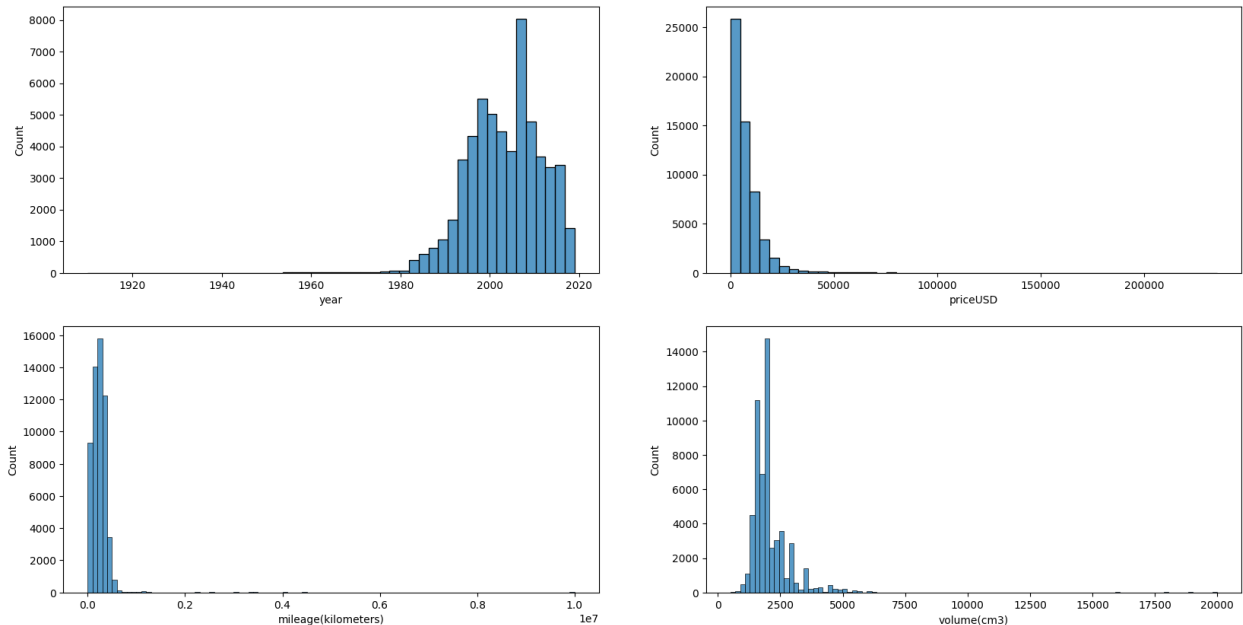
<Axes: xlabel='make_segment', ylabel='count'>
```



Continous variable Distribution

```
fig, ax = plt.subplots(2,2,figsize=(20,10))
sns.histplot(df['year'], ax=ax[0,0],bins=50)
sns.histplot(df['priceUSD'],ax=ax[0,1],bins=50)
sns.histplot(df['mileage(kilometers)'],ax=ax[1,0],bins =100)
sns.histplot(df['volume(cm3)'],ax=ax[1,1],bins=100)
```

<Axes: xlabel='volume(cm3)', ylabel='Count'>



```
df = df[df['year']>1980]
```

Price and Make

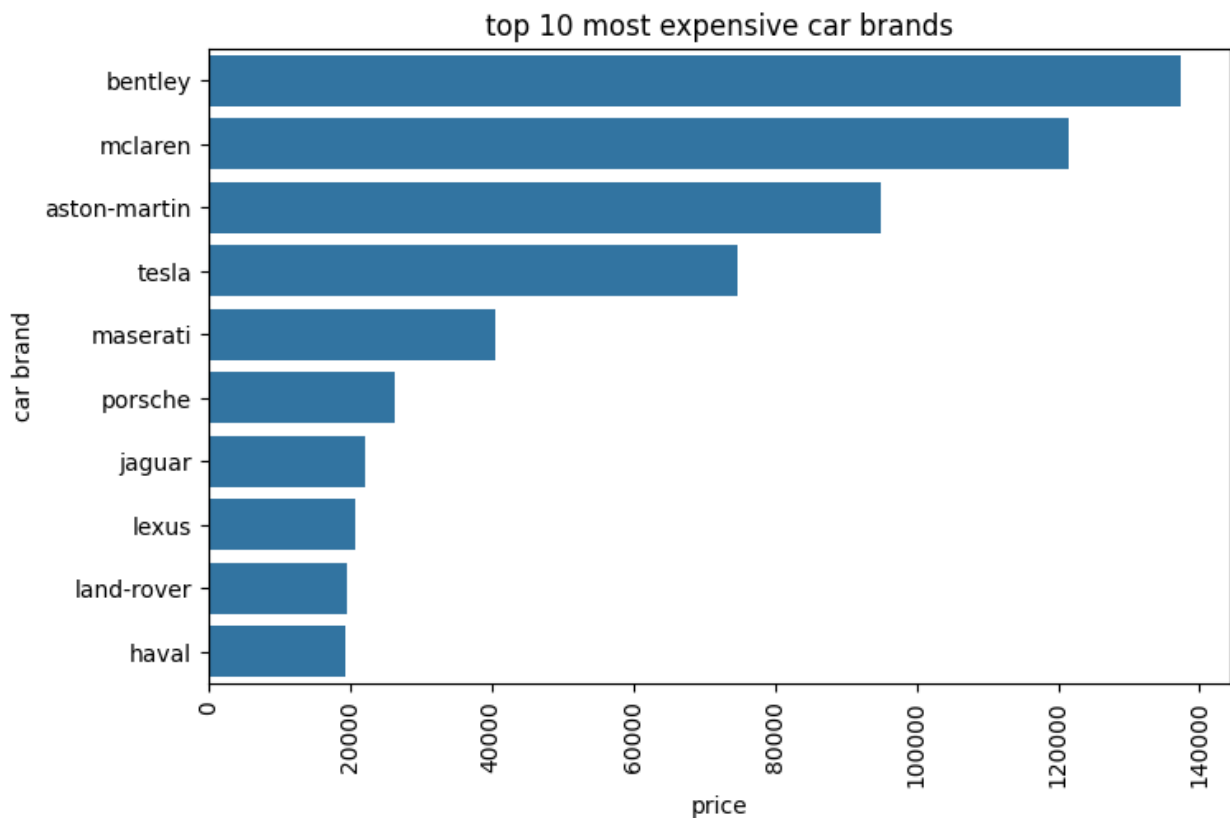
```
demodf = df.groupby('make')['priceUSD'].mean().reset_index()
demof = demodf.sort_values(by='priceUSD',ascending=False).head(10)
```

*# This code is performing the following operations:*  
*#It groups the original dataframe df by the 'make' column.*  
*#For each group (i.e., each unique car make), it calculates the mean of the 'priceUSD' column.*  
*#The result is reset to a new dataframe demodf with two columns: 'make' and 'priceUSD' (which now contains the mean prices).*  
*#The demodf is then sorted by the 'priceUSD' column in descending order (highest price first).*  
*#Finally, it selects only the top 10 rows using .head(10).*  
*#demodf = df.groupby('make')['priceUSD'].mean().reset\_index()*

```
demodf = df.groupby('make')['priceUSD'].mean().reset_index()
demof = demodf.sort_values(by='priceUSD',ascending=False).head(10)
```

*#bar plot*


```
plt.figure(figsize=(8,5))
sns.barplot(y='make', x='priceUSD', data=demof)
plt.xticks(rotation=90)
plt.title("top 10 most expensive car brands")
plt.xlabel("price")
plt.ylabel("car brand")
plt.show()
```



Price and Condition


```
sns.lineplot(x='year', y='priceUSD', data=df, hue = 'condition')
plt.title("price of cars by year and condition")
plt.show()
```






Price and Transmission

```
sns.lineplot(x = 'year', y='priceUSD', data=df, hue = 'transmission')  
plt.title("price of cars and transmission")  
plt.show()
```




Price and Fuel\_type

```
sns.lineplot(x = 'year', y = 'priceUSD', data = df, hue = 'fuel_type')  
plt.title("price of cars and fuel type")  
plt.show()
```




Price and drive\_unit

```
sns.lineplot(x='year', y='priceUSD', data=df, hue='drive_unit')  
plt.title("price and drive unit")  
plt.show()
```



Price and Brand Segment

```
sns.lineplot(x = 'year', y = 'priceUSD', data = df, hue =  
'make_segment')  
plt.title('Price of Cars and Brand Segment')  
plt.show()
```



#### Data Preprocessing Part 2

```
#gives the no of null values totally
df.isnull().sum()

make          0
priceUSD      0
year          0
condition     0
mileage(kilometers)  0
fuel_type     0
volume(cm3)   47
color         0
transmission  0
drive_unit    1874
make_segment  0
dtype: int64

# drops the null values from the columns
df.dropna(inplace=True)

df.isnull().sum()
```

```

make          0
priceUSD      0
year          0
condition     0
mileage(kilometers)  0
fuel_type     0
volume(cm3)   0
color         0
transmission  0
drive_unit    0
make_segment  0
dtype: int64

#to drop null values from the columns(make)
df.drop(columns=['make'], inplace=True)

df.isnull().sum()

priceUSD      0
year          0
condition     0
mileage(kilometers)  0
fuel_type     0
volume(cm3)   0
color         0
transmission  0
drive_unit    0
make_segment  0
dtype: int64

```

Label encoding for object data type

```

from sklearn.preprocessing import LabelEncoder

# columns to encode
cols = ['condition', 'fuel_type', 'transmission', 'color',
        'drive_unit', 'make_segment']

# Label encoding Object is created
le = LabelEncoder()

#label encoding for each column
for col in cols:
    le.fit(df[col])
    df[col] = le.transform(df[col])
    print(col, df[col].unique())

condition [2 1 0]
fuel_type [1 0]
transmission [1 0]


```

```
color [ 3  0 10 11  4  1  7  8  9  5  2 12  6]
drive_unit [1 3 0 2]
make_segment [2 3 5 0 4 6 1]
```

### Correlation Matrix Heatmap

```
#sns.heatmap(): Creates the heatmap using seaborn.
#df.corr(): The correlation matrix to visualize.
#annot=True: Adds numerical annotations to each cell.
#cmap='coolwarm': Sets the color scheme (blue for negative
correlations, red for positive).
plt.figure(figsize=(10,10))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')

<Axes: >
```



## Outlier Removal

Removing outliers is a data preprocessing technique used to eliminate extreme values from a dataset that may skew analysis or model performance.

method for removing Outlier Z-score method: Removes data points that are a certain number of standard deviations away from the mean.

```
# Using Z-score to remove outliers
from scipy import stats
```



```

z = np.abs(stats.zscore(df))

threshold = 3

#columns with outliers
cols = ['year', 'mileage(kilometers)', 'volume(cm3)']

#removing outliers
df = df[(z < 3).all(axis=1)]

```

## Train Test Split

Train-test split is an important technique in machine learning for evaluating model performance.

```

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test =
train_test_split(df.drop(columns=['priceUSD']),df['priceUSD'],
test_size=0.2, random_state=42)

```

## Model Building

### Decision Tree Regressor

**How It Works** The model splits the data into subsets based on feature values, creating a tree-like structure. At each node, it chooses the feature and split point that minimizes the variance in the target variable. The process continues recursively until a stopping criterion is met (e.g., maximum depth reached). Predictions are made by traversing the tree and using the average target value in the leaf node.

### GridSearchCV

GridSearchCV is a powerful tool for hyperparameter tuning in machine learning models. Here are the key points about using GridSearchCV for hyperparameter tuning:

**How GridSearchCV Works** GridSearchCV performs an exhaustive search over a specified parameter grid. It trains the model using every combination of parameters in the grid. Each model is evaluated using cross-validation. The best performing parameter combination is selected.

### Hyperparameter Tuning

Important hyperparameters to consider: **max\_depth**: Controls the maximum depth of the tree  
**min\_samples\_split**: Minimum number of samples required to split a node  
**min\_samples\_leaf**: Minimum number of samples required in a leaf node  
Tuning these parameters helps balance model complexity and performance.

```

from sklearn.model_selection import GridSearchCV

#parameters for grid search
params = {
    'max_depth': [2,4,6,8],

```

```

    'min_samples_split': [2,4,6,8],
    'min_samples_leaf': [1,2,3,4],
    'max_features': ['auto', 'sqrt', 'log2'],
    'random_state': [0,42]
}
# Grid Search Object
grid = GridSearchCV(dtr, param_grid=params, cv=5, verbose=1, n_jobs=-
1)

```

```

#fitting the grid search
grid.fit(x_train, y_train)

```

```

#best parameters
print(grid.best_params_)

```

```

Fitting 5 folds for each of 384 candidates, totalling 1920 fits
{'max_depth': 8, 'max_features': 'sqrt', 'min_samples_leaf': 3,
'min_samples_split': 2, 'random_state': 0}

```

```

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/model_selection/_validation.py:540:

```

```

FitFailedWarning:

```

```

640 fits failed out of a total of 1920.

```

```

The score on these train-test partitions for these parameters will be
set to nan.

```

```

If these failures are not expected, you can try to debug them by
setting error_score='raise'.

```

```

Below are more details about the failures:

```

```

-----

```

```

350 fits failed with the following error:

```

```

Traceback (most recent call last):

```

```

  File

```

```

"/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/model_selection/_validation.py", line 888, in
_fit_and_score

```

```

    estimator.fit(X_train, y_train, **fit_params)

```

```

  File

```

```

"/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/base.py", line 1466, in wrapper

```

```

    estimator._validate_params()

```

```

  File

```

```

"/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/base.py", line 666, in _validate_params

```

```

    validate_parameter_constraints(

```

```

  File

```

```

"/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/utils/_param_validation.py", line 95, in
validate_parameter_constraints

```

```
raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of DecisionTreeRegressor must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'sqrt', 'log2'} or None. Got 'auto' instead.
```

-----  
-----  
290 fits failed with the following error:

Traceback (most recent call last):

File

"/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/model\_selection/\_validation.py", line 888, in
\_fit\_and\_score

estimator.fit(X\_train, y\_train, \*\*fit\_params)

File

"/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/base.py", line 1466, in wrapper

estimator.\_validate\_params()

File

"/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/base.py", line 666, in \_validate\_params

validate\_parameter\_constraints(

File

"/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/utils/\_param\_validation.py", line 95, in

validate\_parameter\_constraints

raise InvalidParameterError(

sklearn.utils.\_param\_validation.InvalidParameterError: The
'max\_features' parameter of DecisionTreeRegressor must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'log2', 'sqrt'} or None. Got 'auto' instead.

warnings.warn(some\_fits\_failed\_message, FitFailedWarning)

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/model\_selection/\_search.py:1102: UserWarning:

One or more of the test scores are non-finite: [ nan nan

nan nan nan nan

nan nan nan nan nan nan

nan nan nan nan nan nan

nan nan nan nan nan nan

nan nan nan nan nan nan

0.30902003 0.293469 0.30902003 0.293469 0.30902003 0.293469

0.30902003 0.293469 0.30902003 0.293469 0.30902003 0.293469

0.30902003 0.293469 0.30902003 0.293469 0.30902003 0.293469

0.30902003 0.293469 0.30902003 0.293469 0.30902003 0.293469

0.30902003 0.293469 0.30902003 0.293469 0.30902003 0.293469

0.30902003 0.293469 0.30902003 0.293469 0.30902003 0.293469

0.30902003 0.293469 0.30902003 0.293469 0.30902003 0.293469

0.30902003	0.293469	0.30902003	0.293469	0.30902003	0.293469
0.30902003	0.293469	0.30902003	0.293469	0.30902003	0.293469
0.30902003	0.293469	0.30902003	0.293469	0.30902003	0.293469
0.30902003	0.293469	0.30902003	0.293469	0.30902003	0.293469
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	0.61141652	0.43704915	0.61141652	0.43704915
0.61141652	0.43704915	0.61141652	0.43704915	0.61141652	0.43704915
0.61141652	0.43704915	0.61141652	0.43704915	0.61141652	0.43704915
0.61141652	0.43704915	0.61141652	0.43704915	0.61141652	0.43704915
0.61141652	0.43704915	0.61141652	0.43704915	0.61141652	0.43704915
0.61141652	0.43704915	0.61141652	0.43704915	0.61141652	0.43704915
0.61141652	0.43704915	0.61141652	0.43704915	0.61141652	0.43704915
0.61141652	0.43704915	0.61141652	0.43704915	0.61141652	0.43704915
0.61141652	0.43704915	0.61141652	0.43704915	0.61141652	0.43704915
0.61141652	0.43704915	0.61141652	0.43704915	0.61141652	0.43704915
0.61141652	0.43704915	0.61141652	0.43704915	0.61141652	0.43704915
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	0.71416767	0.62224628	0.71416767	0.62560501
0.71416767	0.66759186	0.71419186	0.66759186	0.71436534	0.62562559
0.71436534	0.62562559	0.71436534	0.66759186	0.7143242	0.66759186
0.71438433	0.66764212	0.71438433	0.66764212	0.71438433	0.66764212
0.71432582	0.66764212	0.71435735	0.66764125	0.71435735	0.66764125
0.71435735	0.66764125	0.71435735	0.66764125	0.71416767	0.62224628
0.71416767	0.62560501	0.71416767	0.66759186	0.71419186	0.66759186
0.71436534	0.62562559	0.71436534	0.62562559	0.71436534	0.66759186
0.7143242	0.66759186	0.71438433	0.66764212	0.71438433	0.66764212
0.71438433	0.66764212	0.71432582	0.66764212	0.71435735	0.66764125
0.71435735	0.66764125	0.71435735	0.66764125	0.71435735	0.66764125
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	0.77370547	0.78549495	0.78467723	0.77881784
0.77975187	0.76051473	0.77548776	0.7853812	0.77548073	0.77829428
0.77548073	0.77829428	0.78282316	0.77882738	0.77266096	0.76744459
0.78990343	0.78301775	0.78990343	0.78301775	0.78990343	0.78301775
0.77829985	0.78830365	0.77314992	0.78453135	0.77314992	0.78453135
0.77314992	0.78453135	0.77314992	0.78453135	0.77370547	0.78549495
0.78467723	0.77881784	0.77975187	0.76051473	0.77548776	0.7853812
0.77548073	0.77829428	0.77548073	0.77829428	0.78282316	0.77882738

```
0.77266096 0.76744459 0.78990343 0.78301775 0.78990343 0.78301775
0.78990343 0.78301775 0.77829985 0.78830365 0.77314992 0.78453135
0.77314992 0.78453135 0.77314992 0.78453135 0.77314992 0.78453135]
warnings.warn()
```

```
from sklearn.model_selection import GridSearchCV
from sklearn.tree import DecisionTreeRegressor

# Create the base model
dtr = DecisionTreeRegressor()

# Parameters for grid search
params = {
    'max_depth': [2, 4, 6, 8],
    'min_samples_split': [2, 4, 6, 8],
    'min_samples_leaf': [1, 2, 3, 4],
    'max_features': ['sqrt', 'log2', None], # Removed 'auto', added
None
    'random_state': [0, 42]
}

#This creates a GridSearchCV object:
#estimator=dtr: The model to use
#param_grid=params: The hyperparameter grid
#cv=5: Use 5-fold cross-validation
    #Cross-validation is a statistical method used to evaluate machine-
    #learning models by testing them on multiple subsets of data
#verbose=1: Print progress
#n_jobs=-1: Use all available CPU cores
#error_score='raise': Raise errors for debugging

# Grid Search Object
grid = GridSearchCV(
    estimator=dtr,
    param_grid=params,
    cv=5,
    verbose=1,
    n_jobs=-1,
    error_score='raise' # This will raise errors for debugging
)

# Fitting the grid search
grid.fit(x_train, y_train)

# Best parameters
print("Best parameters found:")
print(grid.best_params_)

# Best score
```

```

print("\nBest cross-validation score:
{:.4f}".format(grid.best_score_))

# Best estimator
print("\nBest estimator:")
print(grid.best_estimator_)

Fitting 5 folds for each of 384 candidates, totalling 1920 fits
Best parameters found:
{'max_depth': 8, 'max_features': None, 'min_samples_leaf': 4,
 'min_samples_split': 2, 'random_state': 0}

Best cross-validation score: 0.8573

Best estimator:
DecisionTreeRegressor(max_depth=8, min_samples_leaf=4, random_state=0)

#decision tree regressor with best parameters
dtr = DecisionTreeRegressor(max_depth=8, max_features='sqrt',
min_samples_leaf=4, min_samples_split=2, random_state=0)

#fitting the model
dtr.fit(x_train, y_train)

DecisionTreeRegressor(max_depth=8, max_features='sqrt',
min_samples_leaf=4,
                    random_state=0)

dtr.score(x_train, y_train)

0.7883902435486847

y_pred = dtr.predict(x_test)

```

## Model Evaluation

Model evaluation is a critical process in machine learning that assesses the performance and quality of trained models

### Purpose of Model Evaluation

- 1.Assess model performance and accuracy
- 2.Determine how well a model generalizes to unseen data
- 3.Compare different models or algorithms
- 4.Identify strengths and weaknesses of a model
- 5.Guide model selection and optimization

```

from sklearn.metrics import r2_score, mean_squared_error,
mean_absolute_error
print('R2 Score: ', r2_score(y_test, y_pred))
print('Mean Squared Error: ', mean_squared_error(y_test, y_pred))
print('Mean Absolute Error: ', mean_absolute_error(y_test, y_pred))

```

```
print('Root Mean Squared Error: ', np.sqrt(mean_squared_error(y_test,  
y_pred)))
```

R2 Score: 0.7701494431533182

Mean Squared Error: 7355859.020309841

Mean Absolute Error: 1869.2406549663542

Root Mean Squared Error: 2712.168693188136